

Lecture 3: MPI Basics

Communicators & point-to-point communications (1)

MPI 4 standard: <https://www.mpi-forum.org/docs/mpi-4.0/mpi40-report.pdf>

MPI 3 (version 3.1) standard: <https://www.mpi-forum.org/docs/mpi3.1/mpi31-report.pdf>

<https://www.mpi-forum.org/docs/mpi-2.2/mpi22-report.pdf>

OpenMPI documentation: <https://www.open-mpi.org>

- What would be a “minimal” MPI-based program (which does some communication)? (3 min)

```
MPI_Status status;
```

```
MPI_Init();
```

```
MPI_Sendrecv(&a, 1, MPI_INT, 0, 0, &b, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
```

```
MPI_Finalize();
```

- What is a group, what is a communicator, what is a context?

You already know `MPI_Send` & `MPI_Recv`.

Is their “tag” argument a part of the context? (4 min)

- Group: set of numbered processes
- Communicator: group + context + virtual topology [+ attributes]
= **encapsulator** (think of libraries)
- Context: communicator-specific labels or tags of messages
partition the communication space:
“A message sent in one context cannot be received in another context.”
“Collective operations are independent of point-to-point operations.”
- Message tag: message-specific, not part of communicator context

- About send and receive functions (`MPI_Send`, `MPI_Recv` & other):
Which matching rules do exist for their data type parameters?
Exceptions?
Which parameters can be wildcarded? (5 min)

- Data types must match
 - between send and receive calls (exception: `MPI_PACKED`)
 - between caller declarations and types in MPI calls
(exceptions: `MPI_BYTE` , `MPI_PACKED`)

data counts need not to match
(note: receive count="length of receive buffer", not "length of message")
- Wildcards: `MPI_ANY_SOURCE`, `MPI_ANY_TAG`

- Receive functions:
Which information can be obtained from the **status** parameter? (2 min)
- MPI_SOURCE, MPI_TAG and MPI_ERROR
indirectly: data count by MPI_GET_COUNT
often not relevant → use MPI_STATUS_IGNORE

- Point-to-point communication

Explain the pair of opposites: (10 min)

blocking	– nonblocking
synchronous	– asynchronous
buffered	– unbuffered
local	– nonlocal
one-sided	– two-sided

When are separate completion calls needed?
Give examples for such functions!

- Blocking vs. non-blocking: different in **function returning**
blocking: returns only when success criterion fulfilled
non-blocking: returns immediately
- Asynchronous vs. synchronous: different in **success criterion**
asynchronous: ``buffer is copied"
synchronous: ``data are (beginning to be) received" (= handshake w. receiver)
- Buffered vs. unbuffered: different in **buffer use** (user-provided or system)
buffered: copies data immediately to buffer for later transmission
implies asynchronous
- Local vs. non-local: different in **dependence on other process**
local: returns irrespective of execution within another process
- One-sided vs. two-sided: different in **number of parties** involved
one-sided: only one party (getter or putter) = remote memory access(RMA)
(but 2nd party needs to enable access)
two-sided: two parties - corresponding sender and receiver

- Completion routines: needed for non-blocking calls,
employ **request** parameter
- Examples: MPI_WAIT, MPI_TEST

- Fill the table!

Possibilities: yes/no/depends/not applicable (8 min)

type	blocking	synchronous	buffered	local	remarks
MPI_Bsend	yes	no	yes	yes	
MPI_Ibsend	no	no	yes	yes	
MPI_Ssend	yes	yes	no	no	
MPI_Issend	no	yes	no	yes	
MPI_Send	yes	depends	depends	depends	“standard”
MPI_Isend	no	depends	depends	yes	— ” —
MPI_Rsend	yes	depends	depends	depends	“ready”
MPI_Irsend	no	depends	depends	yes	= receive must be posted
MPI_Sendrecv	yes	depends	depends	no	no deadlocks
MPI_Recv	yes	not applicable	not applicable	no	
MPI_Irecv	no	not applicable	not applicable	yes	