

CS-E4690 – Programming parallel supercomputers D

4th lecture

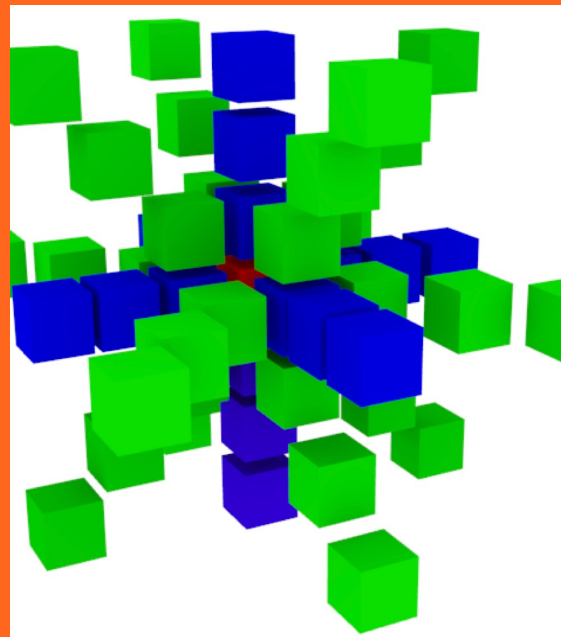
Maarit Korpi-Lagg

maarit.korpi-lagg@aalto.fi

14.11.2023



Aalto University
School of Science



Lecture 4

*Collectives (finalizing basic MPI)
One-sided communication
(entering the advanced domain)*

- **Course practicalities: 5 min**
- **Key concepts recap (40 min) – old and new**
- **Break (max 15 mins)**
- **Example codes cntd. (20 mins)**
- **Exercise sheet tasks tips (20 mins)**
- **Wrap up (poll & feedback; 5 mins)**

Break-down of learning objectives

Lecture1

Introduction to the current HPC landscape

Understanding how this course fits into that

Establishing understanding of the learning outcomes, specifically answering the question: “What are programming during this course?”

Lecture2

Learning basic definitions and taxonomies

Understanding the importance of the “network”

Learning basic performance models

Understanding the concept of a well-performing software in large-scale computing.

Lecture3

Becoming knowledgeable of the modern landscape of distributed memory programming

Understanding why in this course we will concentrate on low-level programming models

Getting acquainted with MPI: basics and synchronous and asynchronous point-to-point communication

Break-down of learning objectives

Lecture4

Learning more about MPI:

One-sided point-to-point communications

Collective communications

Lecture5

Programming MP hybrid architectures

Becoming knowledgeable of the spectrum of options

Understanding efficiency issues

Lecture6

Programming hybrid architectures with accelerators

Acquiring knowledge of CUDA-MPI programming model

Repetition of key concepts - old and new

Aim: to explain the key concept in short

Discuss 1-5 mins in row-wise groups

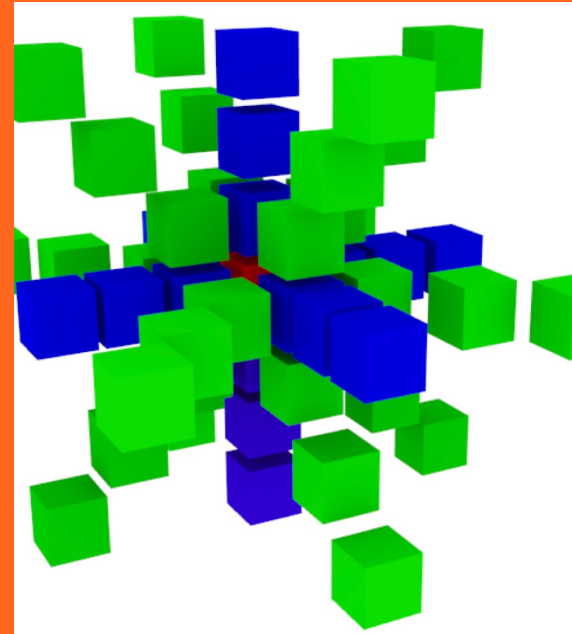
Randomly selected group(s) present(s) and
another one comment(s)

Model answer

Feedback with post-its



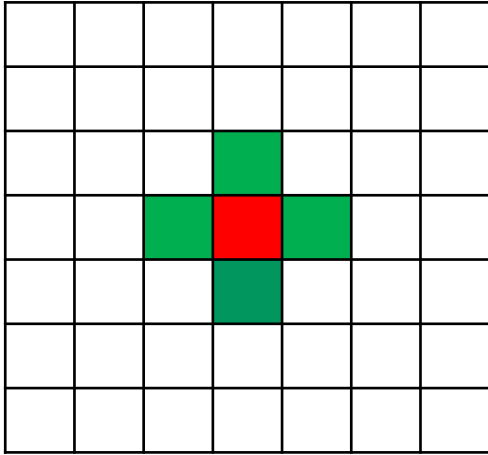
Aalto University
School of Science



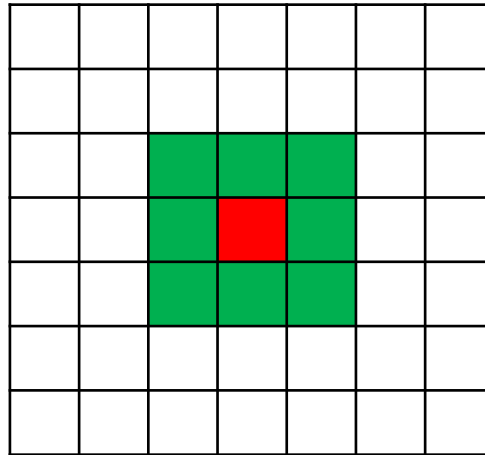
Iterative stencil loop

**Recurring update pattern of array elements
based on their neighbors.**

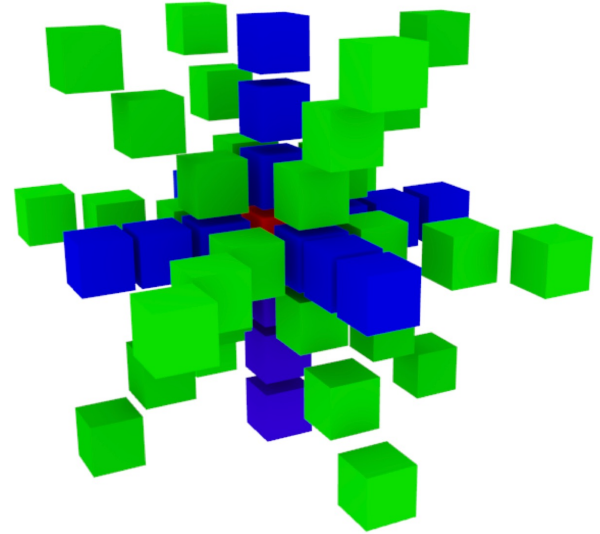
Iterative stencil loop



?



?



?

Draw the stencil of

Sheet 3 ex. 2

$$\frac{\partial c}{\partial x}(x_i, y_j, t_n) \approx \frac{+3c_{i,j}^n - 4c_{i-1,j}^n + c_{i-2,j}^n}{2\Delta x} \quad \text{for} \quad v_x > 0$$
$$\frac{\partial c}{\partial x}(x_i, y_j, t_n) \approx \frac{-c_{i+2,j}^n + 4c_{i+1,j}^n - 3c_{i,j}^n}{2\Delta x} \quad \text{for} \quad v_x < 0,$$

... and the same for y velocity in j direction.

**Two-sided p2p communication;
Examples belonging/not belonging
(at least one per group)?**

Two parties – corresponding sender and receiver

MPI_Send MPI_Irecv
MPI_Ssend
MPI_Isend
MPI_Bsend
MPI_Recv
MPI_Sendrecv
...

MPI_Put
MPI_Get
MPI_Bcast
MPI_Scatter ...
MPI_Accumulate
MPI_Get_accumulate

What does "MPI messages are non-overtaking" mean?

If a sender sends two messages in succession to the same destination, and both match the same receive, then this operation cannot receive the second message if the first one is still pending.

**Does this apply to one-sided
communication?**

It depends. Many MPI_Put and MPI_Get mixed within an epoch do not guarantee the order and can lead to **race conditions.**

Are there ways to avoid race conditions in one-sided communication?

Yes!

Remember to synchronize

OR

**Use MPI_Accumulate with MPI_REPLACE
and MPI_Get_accumulate with MPI_NO_OP
which implement atomic operations for the
same target – origin pairs.**

What is the difference with rooted and all-to-all collectives?

**Results of reductions are collected to
ROOT's receive buffer versus to all ranks'
receive buffers**

MPI_Reduce vs. MPI_Allreduce

Are user defined ops allowed in RMA reductions?

Not yet.

MPI type	meaning	applies to\
MPI_MAX	maximum	integer, floating point
MPI_MIN	minimum	
MPI_SUM	sum	integer, floating point, complex, multilanguage types
MPI_REPLACE	overwrite	
MPI_NO_OP	no change	
MPI_PROD	product	
MPI_LAND	logical and	C integer, logical
MPI_LOR	logical or	
MPI_LXOR	logical xor	
MPI_BAND	bitwise and	integer, byte, multilanguage types
MPI_BOR	bitwise or	
MPI_BXOR	bitwise xor	
MPI_MAXLOC	max value and location	MPI_DOUBLE_INT and such
MPI_MINLOC	min value and location	

**Is there a way to make blocking two-sided
ops safe?**

**Not fully. Using MPI_Sendrecv, MPI_Ssend
and MPI_Bsend may help to write safe
code.**

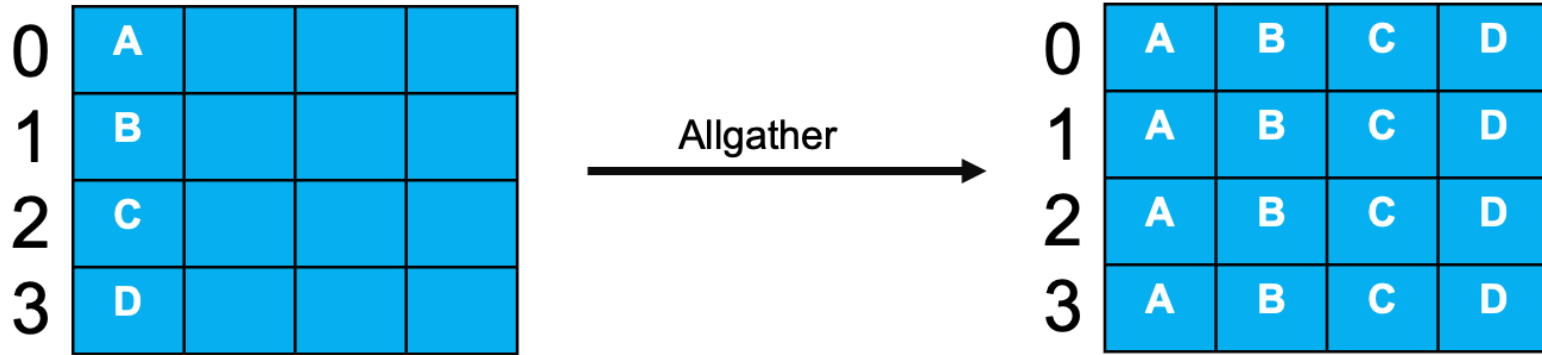
Will blocking two-sided ops give you the optimum performance?

**No, as these functions do not allow for
concurrency in computation and
communication.**

**Will non-blocking two-sided ops give you
the optimum performance?**

It depends: these functions allow for concurrency in computation and communication, but still require some implicit synchronization, buffers, and activity from both sender and receiver.

How could you replace Allgather with simpler MPI collective functions?



First Gather, and then Broadcast

0	A			
1	B			
2	C			
3	D			

Allgather

0	A	B	C	D
1	A	B	C	D
2	A	B	C	D
3	A	B	C	D

What, in terms of matrix operations, is done here?

0	A0	A1	A2	A3
1	B0	B1	B2	B3
2	C0	C1	C2	C3
3	D0	D1	D2	D3

Alltoall →

0	A0	B0	C0	D0
1	A1	B1	C1	D1
2	A2	B2	C2	D2
3	A3	B3	C3	D3

Transpose

0	A0	A1	A2	A3
1	B0	B1	B2	B3
2	C0	C1	C2	C3
3	D0	D1	D2	D3

Alltoall



0	A0	B0	C0	D0
1	A1	B1	C1	D1
2	A2	B2	C2	D2
3	A3	B3	C3	D3



Great
work!

Phew!